

1 Two stochastic epidemic models

We consider two discrete-time, branching-process models for the spread of an infectious disease. Time is measured in days. Let $I_t \geq 0$ denote the number of new infections (incidence) on day t , and let $R_t > 0$ be the effective reproduction number on day t .

The *serial interval weights* $w_s > 0$ ($s = 1, 2, \dots$) satisfy $\sum_{s=1}^{\infty} w_s = 1$ and encode the probability that a secondary infection appears s days after the primary case. Write

$$F_r = \sum_{s=1}^r w_s$$

for the cumulative serial interval distribution at lag r , so $1 - F_r = \sum_{s=r+1}^{\infty} w_s$ is the probability that transmission occurs after day r .

The *dispersion parameter* $k > 0$ controls individual-level variation in infectiousness. Small k corresponds to high heterogeneity (superspreading); as $k \rightarrow \infty$ the negative-binomial offspring distribution converges to a Poisson.

Remark. Throughout, $\text{NB}(\text{mean} = \mu, \text{disp} = k)$ denotes the negative-binomial distribution with probability mass function

$$\mathbb{P}(X = n) = \binom{n+k-1}{n} \left(\frac{k}{k+\mu} \right)^k \left(\frac{\mu}{k+\mu} \right)^n, \quad n = 0, 1, 2, \dots,$$

mean μ , variance $\mu + \mu^2/k$, and probability generating function (pgf) $G(s) = (k / (k + \mu(1-s)))^k$. Similarly, $\text{Gamma}(\alpha, \beta)$ denotes the Gamma distribution with shape α and rate β , so density $f(x) = \beta^\alpha x^{\alpha-1} e^{-\beta x} / \Gamma(\alpha)$, mean α/β , and variance α/β^2 .

2 SSE model (Superspreading Events)

In the *superspreading-events* (SSE) model, the force of infection on day t is the weighted sum of past incidence,

$$\lambda_t = \sum_{s=1}^{\infty} w_s I_{t-s}, \tag{1}$$

and new infections are negative-binomially distributed:

$$I_t \mid \lambda_t \sim \text{NB}(\text{mean} = R_t \lambda_t, \text{disp} = k \lambda_t). \tag{2}$$

The conditional mean of I_t is $R_t \lambda_t$ and its variance is $R_t \lambda_t + R_t^2 \lambda_t / k$, so the NB inflates the Poisson variance by the factor $1 + R_t \lambda_t / k$, capturing superspreading at the population level.

3 SSI model (Superspreading Individuals)

In the *superspreading individuals* (SSI) model, heterogeneity enters at the individual level. Each infected individual has a latent infectivity that is Gamma distributed with mean 1 and variance $1/k$.

Each cohort of I_t infected individuals then has an aggregate infectivity

$$Y_t \mid I_t \sim \text{Gamma}(\text{shape} = k I_t, \text{rate} = k), \tag{3}$$

so $\mathbb{E}[Y_t \mid I_t] = I_t$ and $\text{Var}(Y_t \mid I_t) = I_t / k$. New infections arise as a Poisson process driven by past infectivity,

$$I_t \mid (Y_s)_{s < t} \sim \text{Poisson} \left(R_t \sum_{s=1}^{\infty} w_s Y_{t-s} \right). \tag{4}$$

Remark. Marginalising Y_t from the SSI model recovers a negative-binomial offspring distribution (by the Poisson–Gamma mixture identity; see below), so the two models share the same mean–variance structure at the population level while differing in the *mechanism* generating overdispersion.

4 Extinction probability

Henceforth assume the reproduction number is constant, $R_t = R_0$. We study the probability that the epidemic goes extinct, starting from a single infectious case on day 0.

By branching-process theory, the extinction probability q is the smallest non-negative fixed point of the pgf of the offspring distribution. For both models the offspring distribution is NB(mean = R_0 , disp = k), whose pgf is $G(s) = (k/(k + R_0(1 - s)))^k$. Setting $q = G(q)$ gives

$$q = \left(1 + \frac{R_0}{k}(1 - q)\right)^{-k}. \quad (5)$$

If $R_0 \leq 1$ the unique solution is $q = 1$ (extinction is certain); if $R_0 > 1$ there is a unique solution $q \in (0, 1)$, so the probability of a major outbreak is $1 - q > 0$.

4.1 Extinction probability after observing incidence on days 0 to r

Suppose we observe incidence $I_0, I_1, \dots, I_r$ on the first $r + 1$ days, with $I_0 \geq 1$. We compute the updated extinction probability q_r for the residual epidemic — the probability that the chain triggered by the observed cases eventually dies out.

A standard branching-process argument applies in both models: each remaining offspring of the observed cases, once born, initiates an independent fresh epidemic governed by the original NB(R_0 , k) offspring distribution and goes extinct with the standard probability q (the smallest fixed point of G). The conditioning event constrains only the offspring of the observed cases on days $0, \dots, r$; it does not propagate to later generations. Therefore

$$q_r = G_r(q),$$

where G_r is the pgf of the (conditional) total remaining offspring of the observed cases. The two models differ in the form of G_r .

4.1.1 SSE model

In the SSE model the offspring process decomposes into independent contributions: each of the I_s individuals infected on day s generates a number of offspring at lag ℓ that is NB(mean = $R_0 w_\ell$, disp = $k w_\ell$), independently across individuals, donor days, and lags. By the NB closure property under a common success probability $p = k w_\ell / (k w_\ell + R_0 w_\ell) = k / (k + R_0)$ (independent of ℓ), the population-level NB($R_0 \lambda_t$, $k \lambda_t$) law of I_t is recovered by summing all per-individual-per-lag contributions.

Conditioning on the observed counts $I_0, \dots, I_r$ pins down only those contributions (s, ℓ) with $s + \ell \leq r$. Per-individual-per-lag contributions with $s \in \{0, \dots, r\}$ and $\ell > r - s$ are independent of the observations and retain their original distributions. The total remaining offspring of the observed cases is therefore

$$\sum_{s=0}^r \sum_{i=1}^{I_s} \sum_{\ell=r-s+1}^{\infty} \text{NB}(\text{mean} = R_0 w_\ell, \text{disp} = k w_\ell) = \text{NB}(\text{mean} = R_0 \Lambda, \text{disp} = k \Lambda), \quad (6)$$

where the pooled effective weight is

$$\Lambda = \sum_{s=0}^r I_s (1 - F_{r-s}). \quad (7)$$

With the convention $F_0 = 0$, the $s = r$ term enters with full weight, since none of its offspring have had a chance to appear yet.

The pgf is $G_r(s) = (1 + (R_0/k)(1 - s))^{-k\Lambda} = G(s)^\Lambda$, so

$$q_r = G_r(q) = q^\Lambda = q^{\sum_{s=0}^r I_s (1 - F_{r-s})}. \quad (8)$$

The earlier special case (a single case on day 0 followed by r days without cases) corresponds to $I_0 = 1$, $I_1 = \dots = I_r = 0$, giving $\Lambda = 1 - F_r$ and recovering $q_r = q^{1 - F_r}$. In the same direction, as the zero streak extends ($r \rightarrow \infty$) the contributions $I_s (1 - F_{r-s})$ all vanish and $q_r \rightarrow 1$: extinction approaches certainty.

4.1.2 SSI model

For the SSI model the joint posterior of $(Y_0, \dots, Y_{r-1})$ is in general not available in closed form: as soon as $I_d > 0$ on at least one day $d \in \{1, \dots, r\}$, the day- d rate $R_0 \sum_{s=1}^{d-1} w_{d-s} Y_s$ couples the latents. We therefore present an MCMC estimator of q_r here as the canonical SSI workflow. For the special cases of cases on day 0 only, on day 0 and one later day, and on day 0 and two later days, closed forms are however available; these are deferred to § 6 below.

MCMC posterior. We sample $(Y_0, \dots, Y_{r-1})$ from their joint posterior given $(I_0, \dots, I_r)$ using a Hamiltonian Monte Carlo sampler (implemented in PyMC). Note that Y_r , the aggregate infectivity of the day- r cohort, has not yet contributed to any observed outcome; its posterior equals its prior,

$$Y_r \sim \text{Gamma}(kI_r, k),$$

and is sampled independently.

Remaining offspring from observed cases. Given $(Y_s)_{s=0}^r$, the future infections generated by the observed cases arrive as independent Poisson streams: cohort s contributes at lag $\ell \geq r - s + 1$ at rate $R_0 w_\ell Y_s$. Summing over cohorts and lags gives total remaining offspring

$$I_{\text{rem}} \mid (Y_s)_{s=0}^r \sim \text{Poisson}(R_0 \Lambda), \quad \Lambda = \sum_{s=0}^r Y_s (1 - F_{r-s}), \quad (9)$$

exactly the same pooled-weight formula as in the SSE model.

Extinction probability and MCMC estimator. Each remaining offspring initiates an independent fresh epidemic that goes extinct with the standard probability q . Hence the conditional extinction probability of the residual epidemic, given the latent infectivities, is the pgf of $\text{Poisson}(R_0 \Lambda)$ evaluated at q :

$$q_r \mid (Y_s)_{s=0}^r = e^{R_0 \Lambda (q-1)}. \quad (10)$$

Marginalising over the posterior of (Y_s) and estimating by a Monte Carlo average over M MCMC draws,

$$\hat{q}_r = \frac{1}{M} \sum_{m=1}^M e^{R_0 \Lambda^{(m)} (q-1)}, \quad \widehat{\text{PMO}} = 1 - \hat{q}_r, \quad (11)$$

where $\Lambda^{(m)} = \sum_{s=0}^r Y_s^{(m)} (1 - F_{r-s})$.

Remark (Consistency with the day-0-only analytic formula). When $I_1 = \dots = I_r = 0$ the observations force $Y_1 = \dots = Y_r = 0$, so $\Lambda = Y_0(1 - F_r)$ and the posterior is $Y_0 \sim \text{Gamma}(kI_0, k + R_0 F_r)$. Taking the expectation of $e^{R_0 Y_0 (1 - F_r)(q-1)}$ with respect to this posterior via the Gamma moment-generating function gives

$$q_r = \left(1 - \frac{R_0(1 - F_r)(q-1)}{k + R_0 F_r}\right)^{-kI_0} = \left(1 + \frac{R_0(1 - F_r)}{k + R_0 F_r}(1 - q)\right)^{-kI_0},$$

which matches the closed-form result (22).

5 Probability of major outbreak under model uncertainty

So far we have computed the probability of a major outbreak conditional on a chosen model (SSE or SSI). In practice the underlying mechanism generating overdispersion is unknown, and we may wish to combine the two via Bayesian model averaging. Let $\pi_{\text{SSE}}, \pi_{\text{SSI}} \in [0, 1]$ with $\pi_{\text{SSE}} + \pi_{\text{SSI}} = 1$ be prior probabilities on the two models. Given the observed history $(I_0, \dots, I_r)$, Bayes' theorem yields posterior model probabilities

$$\pi_M^* = \frac{\pi_M L_M}{\pi_{\text{SSE}} L_{\text{SSE}} + \pi_{\text{SSI}} L_{\text{SSI}}}, \quad L_M = \mathbb{P}_M(I_1, \dots, I_r \mid I_0), \quad M \in \{\text{SSE}, \text{SSI}\}. \quad (12)$$

The model-averaged probability of a major outbreak is then

$$\text{PMO}_{\text{avg}} = \pi_{\text{SSE}}^* \cdot \text{PMO}_{\text{SSE}} + \pi_{\text{SSI}}^* \cdot \text{PMO}_{\text{SSI}}, \quad (13)$$

where $\text{PMO}_M = 1 - q_r^M$ as derived above.

5.1 Likelihood of the observed history

SSE model. Conditional on past incidence, the day- t count satisfies $I_t \mid I_{<t} \sim \text{NB}(\text{mean} = R_0 \lambda_t, \text{disp} = k \lambda_t)$ with $\lambda_t = \sum_{s=1}^{t-1} w_{t-s} I_s$, so the joint likelihood factorises into a product of negative-binomial pmfs,

$$L_{\text{SSE}}(I_1, \dots, I_r \mid I_0) = \prod_{t=1}^r \text{NB}(I_t; R_0 \lambda_t, k \lambda_t). \quad (14)$$

This closed form is available for any history.

SSI model, closed-form cases. Closed forms are available in the same three special cases as for q_r (cases on day 0 only, on day 0 and one later day, on day 0 and two later days); these are derived in § 6 below.

SSI model, general histories. The likelihood involves the latent infectivities $(Y_s)_{s=0}^{r-1}$ and is in general not available in closed form. Unfolding the SSI generative dependence $I_0 \rightarrow Y_0 \rightarrow I_1 \rightarrow Y_1 \rightarrow \dots \rightarrow I_r$ as a Markov chain, the joint density of latents and observations factorises as

$$p(Y_{0:r-1}, I_{1:r} \mid I_0) = \prod_{t=0}^{r-1} p_Y(Y_t \mid I_t) \cdot \prod_{t=1}^r p_I(I_t \mid Y_{0:t-1}), \quad (15)$$

where $p_Y(\cdot \mid I_t) = \text{Gamma}(\cdot; k I_t, k)$ is the SSI generative density of Y_t given the cohort size on day t , and $p_I(\cdot \mid Y_{0:t-1}) = \text{Poisson}(\cdot; R_0 \sum_{s=1}^t w_s Y_{t-s})$. Marginalising the latents,

$$L_{\text{SSI}}(I_{1:r} \mid I_0) = \int \prod_{t=0}^{r-1} p_Y(Y_t \mid I_t) \prod_{t=1}^r p_I(I_t \mid Y_{0:t-1}) dY_{0:r-1}. \quad (16)$$

Days s with $I_s = 0$ collapse Y_s to a point mass at 0, and Y_r does not enter L_{SSI} (it would only contribute to $I_{r+1}, I_{r+2}, \dots$), so the integration runs effectively over $\mathcal{T} = \{t \in \{0, \dots, r-1\} : I_t > 0\}$. Three MCMC-based estimators of (16) are derived in § 7.

5.2 Simulation-based estimator

For general histories, where the SSI likelihood is unavailable in closed form, we use a rejection-sampling estimator that simultaneously realises the posterior model weighting. Repeat the following in batches of size b :

1. Draw the per-batch model split as a Binomial: $b_{\text{SSE}} \sim \text{Bin}(b, \pi_{\text{SSE}})$ and $b_{\text{SSI}} = b - b_{\text{SSE}}$.
2. Run b_{SSE} vectorised SSE forward simulations from $t = 0$ seeded with I_0 index cases, retaining only those whose simulated incidence on days $1, \dots, r$ exactly matches the observed $I_1, \dots, I_r$. Likewise for b_{SSI} SSI simulations.
3. For each retained simulation, continue forward until either the single-step incidence reaches the major-outbreak threshold (resolved as “major”) or the most recent $L = \text{len}(w)$ steps are all zero (resolved as “extinct”).

Continue accumulating batches until a target number of resolved-and-matching simulations is reached. The acceptance rate of model M ’s simulations is L_M in expectation; with a Binomial split by prior, the

expected number of accepted M -sims is proportional to $\pi_M L_M$, so the composition of the accepted pool matches the posterior π_M^* . The pooled estimator

$$\widehat{\text{PMO}}_{\text{avg}} = \frac{N_{\text{SSE}}^{\text{maj}} + N_{\text{SSI}}^{\text{maj}}}{N_{\text{SSE}}^{\text{acc}} + N_{\text{SSI}}^{\text{acc}}} = \hat{\pi}_{\text{SSE}}^* \widehat{\text{PMO}}_{\text{SSE}} + \hat{\pi}_{\text{SSI}}^* \widehat{\text{PMO}}_{\text{SSI}} \quad (17)$$

follows immediately, where $\hat{\pi}_M^* = N_M^{\text{acc}} / (N_{\text{SSE}}^{\text{acc}} + N_{\text{SSI}}^{\text{acc}})$ is the implied posterior estimate and $\widehat{\text{PMO}}_M = N_M^{\text{maj}} / N_M^{\text{acc}}$ is the per-model PMO from the accepted M -sims.

6 SSI model: closed forms for histories with cases on a few days

For the SSI model, when non-zero incidence is observed on day 0 and on a small number of later days, the joint posterior of the latent infectivities still factorises (or factorises as a finite mixture) and both the model evidence L_{SSI} and the residual extinction probability q_r admit closed forms. This section treats three such cases in turn: case (i) (§ 6.1), incidence on day 0 only; case (ii) (§ 6.2), incidence on day 0 and one later day; and case (iii) (§ 6.3), incidence on day 0 and two later days.

Strategy. Each derivation follows the same five steps:

1. *Conditional likelihood given the latents.* Multiply the per-day Poisson pmfs to obtain $\mathbb{P}(\text{obs} \mid Y_0, Y_{i_1}, \dots, Y_{i_n})$.
2. *Joint density of latents and observation.* Multiply by the independent Gamma priors $Y_s \sim \text{Gamma}(kI_s, k)$ to get $p(Y_0, Y_{i_1}, \dots, Y_{i_n}, \text{obs})$, with all normalising constants visible.
3. *Likelihood (model evidence).* Integrate the joint density over the latents using the Gamma normalisation $\int_0^\infty Y^{\alpha-1} e^{-\beta Y} dY = \Gamma(\alpha) / \beta^\alpha$ in each variable.
4. *Posterior of the latents.* Divide the joint density by the likelihood to read off the posterior; in the cases below this yields independent Gammas (or a finite mixture of independent-Gamma products in case (iii)).
5. *Residual extinction probability and PMO.* The conditional extinction probability given the latents is $q_r \mid (Y_s) = e^{-R_0(1-q)\Lambda}$ with $\Lambda = \sum_{s=0}^r Y_s(1-F_{r-s})$ (cf. the SSI MCMC derivation in § SSI model above); marginalising over the posterior via the Gamma MGF $\mathbb{E}[e^{-tX}] = (1+t/\beta)^{-\alpha}$ for $X \sim \text{Gamma}(\alpha, \beta)$ gives a closed-form q_r , and the probability of a major outbreak is $\text{PMO} = 1 - q_r$.

A key identity used throughout is that $I_s = 0$ on day s forces $Y_s = 0$ (since $Y_s \mid I_s \sim \text{Gamma}(kI_s, k)$ collapses to a point mass at 0), so among the latents only those tied to non-zero observed days are non-degenerate. Likewise, the SSE likelihood (14) is closed-form for any history, so each subsection below ends with the model-averaged PMO via Bayes' theorem from the per-model likelihoods and PMOs.

6.1 Case (i): $I_0 \geq 1, I_1 = \dots = I_r = 0$

Only Y_0 is non-degenerate; we abbreviate $y = Y_0$.

Step 1: conditional likelihood. Given y , I_d is independently $\text{Poisson}(R_0 w_d y)$ for $d = 1, \dots, r$, so

$$\mathbb{P}(I_1 = \dots = I_r = 0 \mid y) = \prod_{d=1}^r e^{-R_0 w_d y} = e^{-R_0 F_r y}. \quad (18)$$

Step 2: joint density. Multiplying (18) by the prior $y \sim \text{Gamma}(kI_0, k)$,

$$p(y, \text{obs}) = \frac{k^{kI_0}}{\Gamma(kI_0)} y^{kI_0-1} e^{-(k+R_0 F_r)y}. \quad (19)$$

Step 3: likelihood. Integrating y out of (19) via the Gamma normalisation,

$$L_{\text{SSI}}^{(i)} = \frac{k^{kI_0}}{\Gamma(kI_0)} \cdot \frac{\Gamma(kI_0)}{(k + R_0F_r)^{kI_0}} = \left(1 + \frac{R_0F_r}{k}\right)^{-kI_0}. \quad (20)$$

Step 4: posterior. Dividing (19) by (20) every constant cancels except the Gamma rate and shape, giving

$$y \mid \text{obs} \sim \text{Gamma}(kI_0, k + R_0F_r). \quad (21)$$

The posterior rate $k + R_0F_r > k$ reflects the fact that observing no early transmission is evidence that y is small: the index cohort is likely a low-infectivity one.

Step 5: residual extinction probability and PMO. The conditional remaining-offspring rate is $R_0\Lambda = R_0(1 - F_r)y$, so the conditional extinction probability is $q_r \mid y = e^{-R_0(1-q)(1-F_r)y}$. Marginalising over the posterior (21) via the Gamma MGF,

$$q_r = \mathbb{E}\left[e^{-R_0(1-q)(1-F_r)y}\right] = \left(1 + \frac{R_0(1-F_r)(1-q)}{k + R_0F_r}\right)^{-kI_0}, \quad (22)$$

and $\text{PMO} = 1 - q_r$.

Closed-form SSE counterpart. Under the SSE model, $I_t \sim \text{NB}(R_0I_0w_t, kI_0w_t)$ independently across $t = 1, \dots, r$ when $I_1 = \dots = I_r = 0$, so the general likelihood (14) reduces to

$$L_{\text{SSE}} = \prod_{t=1}^r \left(\frac{k}{k + R_0}\right)^{kI_0w_t} = \left(1 + \frac{R_0}{k}\right)^{-kI_0F_r}. \quad (23)$$

The two likelihoods share the same functional shape but with different exponents: the SSE exponent kI_0F_r is reduced by $F_r \leq 1$ relative to the SSI exponent kI_0 , reflecting the fact that the SSE total over days $1, \dots, r$ has dispersion kI_0F_r whereas the SSI total has dispersion kI_0 . Combined with the SSE PMO $1 - q^{I_0(1-F_r)}$ from § SSE model and the SSI PMO (22), the model-averaged PMO is

$$\text{PMO}_{\text{avg}} = \pi_{\text{SSE}}^* (1 - q^{I_0(1-F_r)}) + \pi_{\text{SSI}}^* \left[1 - \left(1 + \frac{R_0(1-F_r)(1-q)}{k + R_0F_r}\right)^{-kI_0}\right]. \quad (24)$$

When $r = 0$, both likelihoods reduce to 1, the posterior reduces to the prior, and the model-averaged PMO is just the prior-weighted average of the two per-model PMOs.

Remark (Consistency with the MCMC formula). When $I_1 = \dots = I_r = 0$, the observations force $Y_1 = \dots = Y_r = 0$, so $\Lambda = Y_0(1 - F_r)$ and the general MCMC posterior of § SSI model collapses to (21). Taking the expectation of $e^{-R_0(1-q)(1-F_r)Y_0}$ with respect to that posterior via the Gamma MGF recovers (22).

6.2 Case (ii): $I_0, I_i \geq 1$ for one later day i , all others zero

We now extend to $I_0 \geq 1$, $I_i \geq 1$ for some $i \in \{1, \dots, r\}$, and $I_d = 0$ for $d \in \{1, \dots, r\} \setminus \{i\}$. Only Y_0 and Y_i are non-degenerate.

Step 1: conditional likelihood. Given (Y_0, Y_i) , the day- d Poisson rate is $R_0w_dY_0$ for $d \in \{1, \dots, i\}$ and $R_0(w_dY_0 + w_{d-i}Y_i)$ for $d \in \{i+1, \dots, r\}$. Multiplying the day- d Poisson pmfs,

$$\mathbb{P}(\text{obs} \mid Y_0, Y_i) = \underbrace{\prod_{d=1}^{i-1} e^{-R_0w_dY_0}}_{= e^{-R_0F_{i-1}Y_0}} \cdot \underbrace{\frac{(R_0w_iY_0)^{I_i}}{I_i!} e^{-R_0w_iY_0}}_{\text{day-}i \text{ pmf}} \cdot \underbrace{\prod_{d=i+1}^r e^{-R_0(w_dY_0 + w_{d-i}Y_i)}}_{= e^{-R_0(F_r - F_i)Y_0} e^{-R_0F_{r-i}Y_i}}.$$

The Y_0 exponentials collapse via $F_{i-1} + w_i + (F_r - F_i) = F_r$, so

$$\mathbb{P}(\text{obs} \mid Y_0, Y_i) = \frac{(R_0w_iY_0)^{I_i}}{I_i!} e^{-R_0F_rY_0} e^{-R_0F_{r-i}Y_i}. \quad (25)$$

Step 2: joint density. Multiplying (25) by the priors $Y_0 \sim \text{Gamma}(kI_0, k)$ and $Y_i \sim \text{Gamma}(kI_i, k)$ (independent),

$$p(Y_0, Y_i, \text{obs}) = \underbrace{\frac{(R_0 w_i)^{I_i}}{I_i!} \cdot \frac{k^{kI_0 + kI_i}}{\Gamma(kI_0) \Gamma(kI_i)}}_{\text{constants}} \times \underbrace{Y_0^{kI_0 + I_i - 1} e^{-(k + R_0 F_r) Y_0}}_{\text{Gamma kernel in } Y_0} \cdot \underbrace{Y_i^{kI_i - 1} e^{-(k + R_0 F_{r-i}) Y_i}}_{\text{Gamma kernel in } Y_i}. \quad (26)$$

Step 3: likelihood. Integrating each Gamma kernel of (26) via the Gamma normalisation gives

$$L_{\text{SSI}}^{(\text{ii})} = \frac{(R_0 w_i)^{I_i}}{I_i!} \frac{\Gamma(kI_0 + I_i)}{\Gamma(kI_0)} \frac{k^{kI_0}}{(k + R_0 F_r)^{kI_0 + I_i}} \left(\frac{k}{k + R_0 F_{r-i}} \right)^{kI_i}. \quad (27)$$

Setting $I_i = 0$ collapses the Γ -ratio and the second parenthesis to 1, recovering (20).

Step 4: posterior. Dividing (26) by (27), every constant cancels and the two Gamma kernels normalise independently:

$$Y_0 \mid \text{obs} \sim \text{Gamma}(kI_0 + I_i, k + R_0 F_r), \quad Y_i \mid \text{obs} \sim \text{Gamma}(kI_i, k + R_0 F_{r-i}). \quad (28)$$

The shape of Y_0 is enriched by I_i : each observed offspring of the index cohort acts as a Bayesian data point tilting Y_0 toward higher values. The Y_i posterior is the day-0-only update (21) applied to the day- i cohort over its residual horizon $r - i$; setting $i = r$ gives $F_{r-i} = F_0 = 0$, so Y_r collapses to its prior, since the day- r cohort has not had a chance to contribute.

Step 5: residual extinction probability and PMO. The remaining-offspring rate is $R_0 \Lambda$ with $\Lambda = Y_0(1 - F_r) + Y_i(1 - F_{r-i})$, so the conditional extinction probability is $q_r \mid (Y_0, Y_i) = e^{-R_0(1-q)\Lambda}$. Marginalising over the two independent Gamma factors of (28) via the Gamma MGF gives

$$q_r = \left(1 + \frac{R_0(1 - F_r)(1 - q)}{k + R_0 F_r} \right)^{-(kI_0 + I_i)} \left(1 + \frac{R_0(1 - F_{r-i})(1 - q)}{k + R_0 F_{r-i}} \right)^{-kI_i}, \quad (29)$$

and $\text{PMO} = 1 - q_r$.

Remark (Reduction to case (i)). Setting $I_i = 0$ in (29) collapses the second factor to 1 and the first factor's exponent to $-kI_0$, recovering (22); likewise (27) reduces to (20). As the observation horizon extends ($r \rightarrow \infty$ with i fixed), both factors of (29) tend to 1 and $q_r \rightarrow 1$.

6.3 Case (iii): $I_0, I_i, I_j \geq 1$ for two later days $i < j$, all others zero

Finally, take $I_0 \geq 1$, $I_i \geq 1$, $I_j \geq 1$ for some $1 \leq i < j \leq r$, and $I_d = 0$ for $d \in \{1, \dots, r\} \setminus \{i, j\}$. Now only Y_0, Y_i, Y_j are non-degenerate.

Step 1: conditional likelihood. The conditional likelihood acquires an extra factor for the day- j Poisson pmf, whose rate $R_0(w_j Y_0 + w_{j-i} Y_i)$ couples Y_0 and Y_i . Repeating the telescoping argument of case (ii) for each latent separately,

$$\mathbb{P}(\text{obs} \mid Y_0, Y_i, Y_j) = \frac{(R_0 w_i Y_0)^{I_i}}{I_i!} \frac{(R_0(w_j Y_0 + w_{j-i} Y_i))^{I_j}}{I_j!} e^{-R_0 F_r Y_0} e^{-R_0 F_{r-i} Y_i} e^{-R_0 F_{r-j} Y_j}. \quad (30)$$

Expanding the day- j rate via the binomial theorem,

$$(w_j Y_0 + w_{j-i} Y_i)^{I_j} = \sum_{m=0}^{I_j} \binom{I_j}{m} w_j^{I_j-m} w_{j-i}^m Y_0^{I_j-m} Y_i^m. \quad (31)$$

Remark (Latent-partition interpretation). Equivalently, by Poisson superposition we may regard I_j as the sum of two independent Poissons — one with mean $R_0 w_j Y_0$ counting day- j infections that originated from the day-0 cohort, and one with mean $R_0 w_{j-i} Y_i$ counting those from the day- i cohort. Each term m in (31) is the joint probability that exactly $I_j - m$ originate from Y_0 and m from Y_i .

Step 2: joint density. Multiplying (30) by the three independent priors $\text{Gamma}(kI_0, k)$, $\text{Gamma}(kI_i, k)$, $\text{Gamma}(kI_j, k)$ and inserting (31),

$$p(Y_0, Y_i, Y_j, \text{obs}) = \frac{(R_0 w_i)^{I_i} R_0^{I_j}}{I_i! I_j!} \frac{k^{kI_0+kI_i+kI_j}}{\Gamma(kI_0) \Gamma(kI_i) \Gamma(kI_j)} Y_j^{kI_j-1} e^{-(k+R_0 F_{r-j})Y_j} \\ \times \sum_{m=0}^{I_j} \binom{I_j}{m} w_j^{I_j-m} w_{j-i}^m Y_0^{kI_0+I_i+I_j-m-1} e^{-(k+R_0 F_r)Y_0} Y_i^{kI_i+m-1} e^{-(k+R_0 F_{r-i})Y_i}. \quad (32)$$

This is a constant times a Gamma kernel in Y_j (the same in every mixture component) times a finite mixture (over $m \in \{0, \dots, I_j\}$) of Gamma-product kernels in (Y_0, Y_i) .

Step 3: likelihood. Integrating (32) via the Gamma normalisation in each variable, define the per-mixture-component constants

$$c_m := \binom{I_j}{m} w_j^{I_j-m} w_{j-i}^m \frac{\Gamma(kI_0 + I_i + I_j - m)}{(k + R_0 F_r)^{kI_0+I_i+I_j-m}} \frac{\Gamma(kI_i + m)}{(k + R_0 F_{r-i})^{kI_i+m}}. \quad (33)$$

Cancelling the $\Gamma(kI_j)$ factors and grouping k^{kI_j} with the Y_j -rate gives

$$L_{\text{SSI}}^{(\text{iii})} = \frac{(R_0 w_i)^{I_i} R_0^{I_j}}{I_i! I_j!} \frac{k^{kI_0+kI_i}}{\Gamma(kI_0) \Gamma(kI_i)} \left(\frac{k}{k + R_0 F_{r-j}} \right)^{kI_j} \sum_{m=0}^{I_j} c_m. \quad (34)$$

Setting $I_j = 0$ collapses the sum to its $m = 0$ term and the $(k/(k + R_0 F_{r-j}))^{kI_j}$ factor to 1, recovering (27).

Step 4: posterior. Dividing (32) by (34): the Y_j -Gamma rate-and-shape pair appears in both numerator and denominator and gives the $\text{Gamma}(Y_j | kI_j, k + R_0 F_{r-j})$ density. Within the sum, the constants outside cancel; what remains in the m th term is c_m (numerator) divided by $\sum_{m'} c_{m'}$ (denominator), times the two unnormalised (Y_0, Y_i) Gamma kernels each with their proper normalising constants supplied directly from the c_m in the denominator. The result is a finite mixture of independent Gamma products,

$$p(Y_0, Y_i, Y_j | \text{obs}) = \text{Gamma}(Y_j | kI_j, k + R_0 F_{r-j}) \\ \times \sum_{m=0}^{I_j} \rho_m \text{Gamma}(Y_0 | kI_0 + I_i + I_j - m, k + R_0 F_r) \text{Gamma}(Y_i | kI_i + m, k + R_0 F_{r-i}), \quad (35)$$

with normalised mixture weights

$$\rho_m = \frac{c_m}{\sum_{m'=0}^{I_j} c_{m'}}. \quad (36)$$

The mixture index m tracks how many of the I_j day- j infections are attributed to Y_i versus Y_0 ; the weights ρ_m encode which partition is most consistent with the prior infectivity distributions, the serial-interval weights w_j , w_{j-i} , and the dispersion k .

Step 5: residual extinction probability and PMO. The remaining-offspring rate is now $R_0 \Lambda = R_0[Y_0(1 - F_r) + Y_i(1 - F_{r-i}) + Y_j(1 - F_{r-j})]$, so the conditional extinction probability is $q_r | (Y_0, Y_i, Y_j) = e^{-R_0(1-q)\Lambda}$. Marginalising over the posterior (35), Y_j 's expectation factors out (its density is independent of m); within each mixture component the (Y_0, Y_i) -densities are independent Gammas whose MGFs evaluate term-by-term, giving

$$q_r = \beta_j^{-kI_j} \sum_{m=0}^{I_j} \rho_m \beta_0^{-(kI_0+I_i+I_j-m)} \beta_i^{-(kI_i+m)}, \quad (37)$$

where, for $a \in \{0, i, j\}$ (with the convention $F_{r-0} = F_r$),

$$\beta_a = 1 + \frac{R_0(1-q)(1-F_{r-a})}{k + R_0 F_{r-a}}.$$

Remark (Reductions to simpler cases). When $I_j \rightarrow 0$ the sum collapses to its $m = 0$ term ($\rho_0 = 1$) and $\beta_j^{-kI_j} \rightarrow 1$: (37) reduces to (29). When also $I_i \rightarrow 0$, the exponent on β_i in (29) vanishes and the exponent on β_0 reduces from $-(kI_0 + I_i)$ to $-kI_0$, recovering (22). Similarly (34) reduces to (27) and then to (20).

Computational note. The c_m involve Γ factors of size up to $kI_0 + I_i + I_j$, so for moderate I_j a log-domain computation ($\log \Gamma$, logsumexp) is recommended. The number of sum terms is $I_j + 1$, typically small for the histories of interest.

7 MCMC-based estimators of the SSI marginal likelihood

We describe three estimators of the integral (16). Writing $\pi(Y) = \prod_{t \in \mathcal{T}} p_Y(Y_t | I_t)$ for the integrand's prior factor restricted to the active days $\mathcal{T} = \{t \in \{0, \dots, r-1\} : I_t > 0\}$ (a product of independent Gammas, treating the observed cohort sizes I_t as fixed shape parameters), and $\mathcal{L}(Y) = \prod_{t=1}^r p_I(I_t | Y_{0:t-1})$ for the integrand's likelihood factor, the target is

$$L_{\text{SSI}} = \int \mathcal{L}(Y) \pi(Y) dY = \mathbb{E}_{Y \sim \pi}[\mathcal{L}(Y)]. \quad (38)$$

The unnormalised posterior (the joint integrand) is $\tilde{p}(Y) = \mathcal{L}(Y) \pi(Y)$, and the posterior density of the latents given the observations is $p_{\text{post}}(Y) = \tilde{p}(Y)/L_{\text{SSI}}$. The MCMC sampler of § 4.1.2 produces draws $Y^{p.(j)} \sim p_{\text{post}}$ for $j = 1, \dots, N_1$.

7.1 Naive Monte Carlo

Equation (38) is itself an expectation under the prior π , so the simplest estimator just averages the conditional likelihood at independent draws from π :

$$\hat{L}_{\text{SSI}}^{\text{naive}} = \frac{1}{N} \sum_{n=1}^N \mathcal{L}(Y^{(n)}), \quad Y^{(n)} \stackrel{\text{iid}}{\sim} \pi. \quad (39)$$

No MCMC trace is needed: drawing $Y^{(n)}$ requires only independent $\text{Gamma}(kI_t, k)$ samples for each $t \in \mathcal{T}$. In log space the implementation is

$$\log \hat{L}_{\text{SSI}}^{\text{naive}} = \text{logsumexp}_{n=1, \dots, N} \log \mathcal{L}(Y^{(n)}) - \log N.$$

The estimator is unbiased. Its relative variance is $\text{Var}_{\pi}[\mathcal{L}]/(NL_{\text{SSI}}^2)$, equivalently $(\chi^2(p_{\text{post}} \| \pi) + 1)/N - 1/N$, where $\chi^2(p \| q) = \int (p/q - 1)^2 q$ is the χ^2 divergence. The χ^2 between prior and posterior is small whenever the data are weakly informative about the latents — and remarkably, in the SSI model this remains true even on moderately informative histories, because the prior variance I_t/k scales automatically with the cohort size I_t and so retains overlap with the posterior. The estimator becomes unreliable only when the data are very informative (e.g. long histories with substantial incidence at small k).

7.2 Importance sampling

Importance sampling generalises the naive estimator by sampling from a proposal g chosen to better cover the posterior:

$$\hat{L}_{\text{SSI}}^{\text{IS}} = \frac{1}{N} \sum_{n=1}^N \frac{\tilde{p}(Y^{(n)})}{g(Y^{(n)})}, \quad Y^{(n)} \stackrel{\text{iid}}{\sim} g. \quad (40)$$

The estimator is unbiased for any proposal g with $g(Y) > 0$ wherever $\tilde{p}(Y) > 0$; its relative variance is $\chi^2(p_{\text{post}} \| g)/N$, minimised by taking g as close to p_{post} as possible.

We take

$$g(Y) = \prod_{t \in \mathcal{T}} \text{Gamma}(Y_t; \alpha_t, \beta_t), \quad (41)$$

i.e. a product of independent Gammas in the same family as the prior, with shape and rate parameters fitted by method of moments to the per-day marginals of the MCMC posterior trace $\{Y^{p,(j)}\}_{j=1}^{N_1}$:

$$\alpha_t = \frac{\bar{Y}_t^2}{S_t^2}, \quad \beta_t = \frac{\bar{Y}_t}{S_t^2}, \quad t \in \mathcal{T}, \quad (42)$$

where $\bar{Y}_t$ and S_t^2 are the sample mean and (unbiased) sample variance of $\{Y_t^{p,(j)}\}_j$. Matching the proposal family to the prior keeps importance weights bounded; matching its marginals to the MCMC posterior makes it close to p_{post} marginally.

Remark (Caveat: posterior correlations). Equation (41) is a product proposal, while the joint posterior of $Y_{\mathcal{T}}$ may have non-trivial correlations arising from the shared force-of-infection terms. When such correlations are strong, $\chi^2(p_{\text{post}} \| g)$ does not shrink to zero in the large-data limit and the IS estimator's variance remains bounded below by that mismatch. Bridge sampling (§ 7.3) is more forgiving of proposal mismatch.

7.3 Bridge sampling

Bridge sampling uses draws from *both* the proposal g and the posterior p_{post} , weighted by a bridge function α . The starting identity (Meng & Wong, 1996) is the algebraic equality

$$L_{\text{SSI}} = \frac{\mathbb{E}_{Y \sim g}[\tilde{p}(Y) \alpha(Y)]}{\mathbb{E}_{Y \sim p_{\text{post}}}[g(Y) \alpha(Y)]}, \quad (43)$$

valid for any α such that $0 < \int g(Y) \tilde{p}(Y) \alpha(Y) dY < \infty$. The identity is verified by writing the numerator as $\int \tilde{p} g \alpha dY$ and the denominator as $L_{\text{SSI}}^{-1} \int \tilde{p} g \alpha dY$. Replacing the expectations by Monte Carlo averages over $\{Y^{q,(i)}\}_{i=1}^{N_2} \stackrel{\text{iid}}{\sim} g$ and $\{Y^{p,(j)}\}_{j=1}^{N_1} \sim p_{\text{post}}$ (the MCMC trace) gives the bridge-sampling estimator.

The asymptotic variance of $\hat{L}_{\text{SSI}}$ depends on the choice of α . Minimising it (Meng & Wong, 1996) yields the optimal bridge

$$\alpha^*(Y) = \frac{C}{s_1 g(Y) + s_2 \tilde{p}(Y)/L_{\text{SSI}}}, \quad (44)$$

where $s_1 = N_1/(N_1 + N_2)$, $s_2 = N_2/(N_1 + N_2)$, and C is any non-zero constant (any rescaling of α cancels between numerator and denominator of (43)). The optimal bridge depends on the unknown L_{SSI} ; substituting (44) into (43) and replacing the expectations by Monte Carlo averages produces the fixed-point identity

$$L_{\text{SSI}} = \frac{\frac{1}{N_2} \sum_{i=1}^{N_2} \frac{\tilde{p}(Y^{q,(i)})}{s_1 g(Y^{q,(i)}) + s_2 \tilde{p}(Y^{q,(i)})/L_{\text{SSI}}}}{\frac{1}{N_1} \sum_{j=1}^{N_1} \frac{g(Y^{p,(j)})}{s_1 g(Y^{p,(j)}) + s_2 \tilde{p}(Y^{p,(j)})/L_{\text{SSI}}}}}. \quad (45)$$

Solving (45) for L_{SSI} requires iterating the right-hand side from an initial guess. Bridge sampling has lower (or equal) asymptotic variance than importance sampling with the same proposal, and the variance bound is governed by a symmetrised divergence rather than $\chi^2(p_{\text{post}} \| g)$, so it is significantly more robust to proposal-posterior mismatch.

In practice we work entirely in log space for numerical stability. Define the log importance ratios at the two sample sets,

$$\ell_q^{(i)} = \log \tilde{p}(Y^{q,(i)}) - \log g(Y^{q,(i)}), \quad \ell_p^{(j)} = \log \tilde{p}(Y^{p,(j)}) - \log g(Y^{p,(j)}).$$

Multiplying numerator and denominator of (45) by $1/g(Y^{q,(i)})$ and $1/g(Y^{p,(j)})$ respectively and taking logs gives the log-space iteration

$$\begin{aligned} \log L^{(t+1)} = & \log \text{sumexp}_{i=1,\dots,N_2} \left[\ell_q^{(i)} - \text{lse}(\log s_1, \log s_2 + \ell_q^{(i)} - \log L^{(t)}) \right] - \log N_2 \\ & - \log \text{sumexp}_{j=1,\dots,N_1} \left[-\text{lse}(\log s_1, \log s_2 + \ell_p^{(j)} - \log L^{(t)}) \right] + \log N_1, \end{aligned} \quad (46)$$

where $\text{lse}(a, b) = \log(e^a + e^b)$. We use the same Gamma proposal g as in § 7.2, initialise $\log L^{(0)}$ at the median of $\{\ell_q^{(i)}\}_i$, and iterate (46) until $\log L^{(t)}$ stops changing appreciably.

8 Uncertainty in the parameters R_0 and k

The preceding sections have treated R_0 and k as known. In practice they are estimated from the data; we equip them with a prior $\pi(R_0, k)$ and ask for the resulting marginal model evidence and probability of a major outbreak. We collect the parameters into $\theta = (R_0, k)$ and assume the prior factorises as $\pi(\theta) = \pi(R_0) \pi(k)$ with each factor either a Gamma or LogNormal density; a fixed parameter corresponds to its prior factor being a point mass, which drops out of the integrals below without affecting the structure of any estimator.

Conditional on θ , the SSE evidence is the closed-form negative-binomial product (14), and the SSI evidence is the latent integral (16); we re-label these conditional evidences $L_M(\theta)$ for $M \in \{\text{SSE}, \text{SSI}\}$ to make the parameter dependence explicit. The corresponding parameter-known PMOs of §§ 4.1.1–4.1.2 we re-label $\text{PMO}_M(\theta)$. When θ is uncertain both quantities become marginals over the prior.

8.1 Marginal evidence and model-averaged PMO

Define the prior-marginal evidence

$$\bar{L}_M = \int L_M(\theta) \pi(\theta) d\theta, \quad M \in \{\text{SSE}, \text{SSI}\}, \quad (47)$$

and the corresponding marginal probability of a major outbreak,

$$\overline{\text{PMO}}_M = \int \text{PMO}_M(\theta) p(\theta | I_{1:r}, M) d\theta = \frac{1}{\bar{L}_M} \int \text{PMO}_M(\theta) L_M(\theta) \pi(\theta) d\theta, \quad (48)$$

the posterior expectation of $\text{PMO}_M(\theta)$ under the parameter posterior of model M . In the limit of a degenerate prior at θ_0 , (47) reduces to $L_M(\theta_0)$ and (48) reduces to $\text{PMO}_M(\theta_0)$, recovering the parameter-known quantities of the preceding sections.

The Bayesian model-averaging formulae of § 5 carry over unchanged. The posterior model probabilities are $\pi_M^* \propto \pi_M \bar{L}_M$ and the model-averaged probability of a major outbreak is $\pi_{\text{SSE}}^* \overline{\text{PMO}}_{\text{SSE}} + \pi_{\text{SSI}}^* \overline{\text{PMO}}_{\text{SSI}}$.

8.2 Estimators on the augmented latent vector

The three estimators of § 7 carry over to (47) after augmenting the latent vector with the uncertain parameters. Write $\theta = (\theta_u, \theta_f)$ where θ_u collects the components of θ carrying a non-degenerate prior factor and θ_f collects the remaining (fixed) components. The augmented latent vector is

$$z = (\theta_u, Y_{\mathcal{T}}),$$

where $Y_{\mathcal{T}}$ is the SSI latent infectivities on the active days $\mathcal{T} = \{t \in \{0, \dots, r-1\} : I_t > 0\}$; for the SSE model this block is empty and $z = \theta_u$. The unnormalised target on z is

$$\tilde{p}(z) = \mathcal{L}(z; \theta_f) \pi_Y(Y_{\mathcal{T}} | k) \pi(\theta_u), \quad (49)$$

where $\mathcal{L}(z; \theta_f) = p(I_{1:r} | z, \theta_f)$ is the observation likelihood (the SSE NB product (14) for SSE; the latent-conditional Poisson product $\prod_{t=1}^r p_I(I_t | Y_{0:t-1})$ of (15) for SSI) and π_Y is the SSI latent prior of § 7 (absent in the SSE case). By construction

$$\bar{L}_M = \int \tilde{p}(z) dz, \quad (50)$$

so $\bar{L}_M$ is the normalising constant of the augmented target and $p_{\text{post}}(z) = \tilde{p}(z)/\bar{L}_M$ is the augmented posterior. When θ_u is empty, (49) collapses to the latent-only target $\tilde{p}(Y)$ of (38) and the development below reduces to that of § 7.

Naive Monte Carlo. Draw $\theta_u^{(n)} \sim \pi$ and, for the SSI model, $Y_{\mathcal{T}}^{(n)} \mid k^{(n)} \sim \pi_Y(\cdot \mid k^{(n)})$, jointly forming $z^{(n)}$. The estimator

$$\hat{\bar{L}}_M^{\text{naive}} = \frac{1}{N} \sum_{n=1}^N \mathcal{L}(z^{(n)}; \theta_f) \quad (51)$$

is unbiased. For the SSE model the latent block is empty and $\mathcal{L}(z; \theta_f)$ is the closed-form $L_{\text{SSE}}(\theta^{(n)})$ of (14); no MCMC is required. Naive MC inherits the variance pathology of § 7.1: it is reliable when the prior and the augmented posterior overlap broadly, and degrades on informative histories where the posterior concentrates in a region of low prior mass.

Importance sampling. Run HMC against $\tilde{p}$ of (49) to obtain posterior draws of the augmented vector z . Fit an independent-Gamma proposal

$$g(z) = \prod_j \text{Gamma}(z_j; \alpha_j, \beta_j) \quad (52)$$

by method of moments (42) to each marginal of the trace (the index j now ranges over the components of θ_u as well as the active days $\mathcal{T}$). The importance-sampling estimator

$$\hat{\bar{L}}_M^{\text{IS}} = \frac{1}{N} \sum_{n=1}^N \frac{\tilde{p}(z^{(n)})}{g(z^{(n)})}, \quad z^{(n)} \stackrel{\text{iid}}{\sim} g, \quad (53)$$

is the literal extension of (40); the variance discussion of § 7.2 applies verbatim with $\chi^2(p_{\text{post}} \parallel g)$ now measured on the augmented space. The Gamma proposal again matches the support of every component of z : the latent infectivities are Gamma a priori, and both Gamma and LogNormal parameter priors are supported on the positive reals.

Bridge sampling. With $\{z^{p,(j)}\}_{j=1}^{N_1}$ from the HMC trace and $\{z^{q,(i)}\}_{i=1}^{N_2}$ from the proposal (52), define the augmented log importance ratios

$$\ell_q^{(i)} = \log \tilde{p}(z^{q,(i)}) - \log g(z^{q,(i)}), \quad \ell_p^{(j)} = \log \tilde{p}(z^{p,(j)}) - \log g(z^{p,(j)}).$$

The Meng–Wong identity (43) and its optimal bridge α^* are unchanged in form (with $\bar{L}_M$ replacing L_{SSI} and the augmented variables in place of Y), and the log-space fixed-point iteration (46) applies verbatim.

8.3 Simulation-based rejection scheme

The rejection sampler of § 5.2 extends without modification: for each simulated trajectory under model M , draw a fresh $\theta^{(i)} \sim \pi$, simulate the SSE or SSI generative process forward with parameters $\theta^{(i)}$, and retain only the trajectories whose first r cohorts match the observed $I_{1:r}$. The acceptance probability of a single M -trajectory is $\int p(I_{1:r} \mid \theta, M) \pi(\theta) d\theta = \bar{L}_M$, so the empirical fraction of accepted trajectories that go on to a major outbreak is consistent for $\bar{\text{PMO}}_M$ of (48). Pooling SSE and SSI trajectories with per-batch model split $\text{Bin}(b, \pi_{\text{SSE}})$ realises the posterior model partition with expected fractions $\pi_M \bar{L}_M$, so the composition of the accepted pool delivers the model-averaged PMO exactly as in § 5.2.